

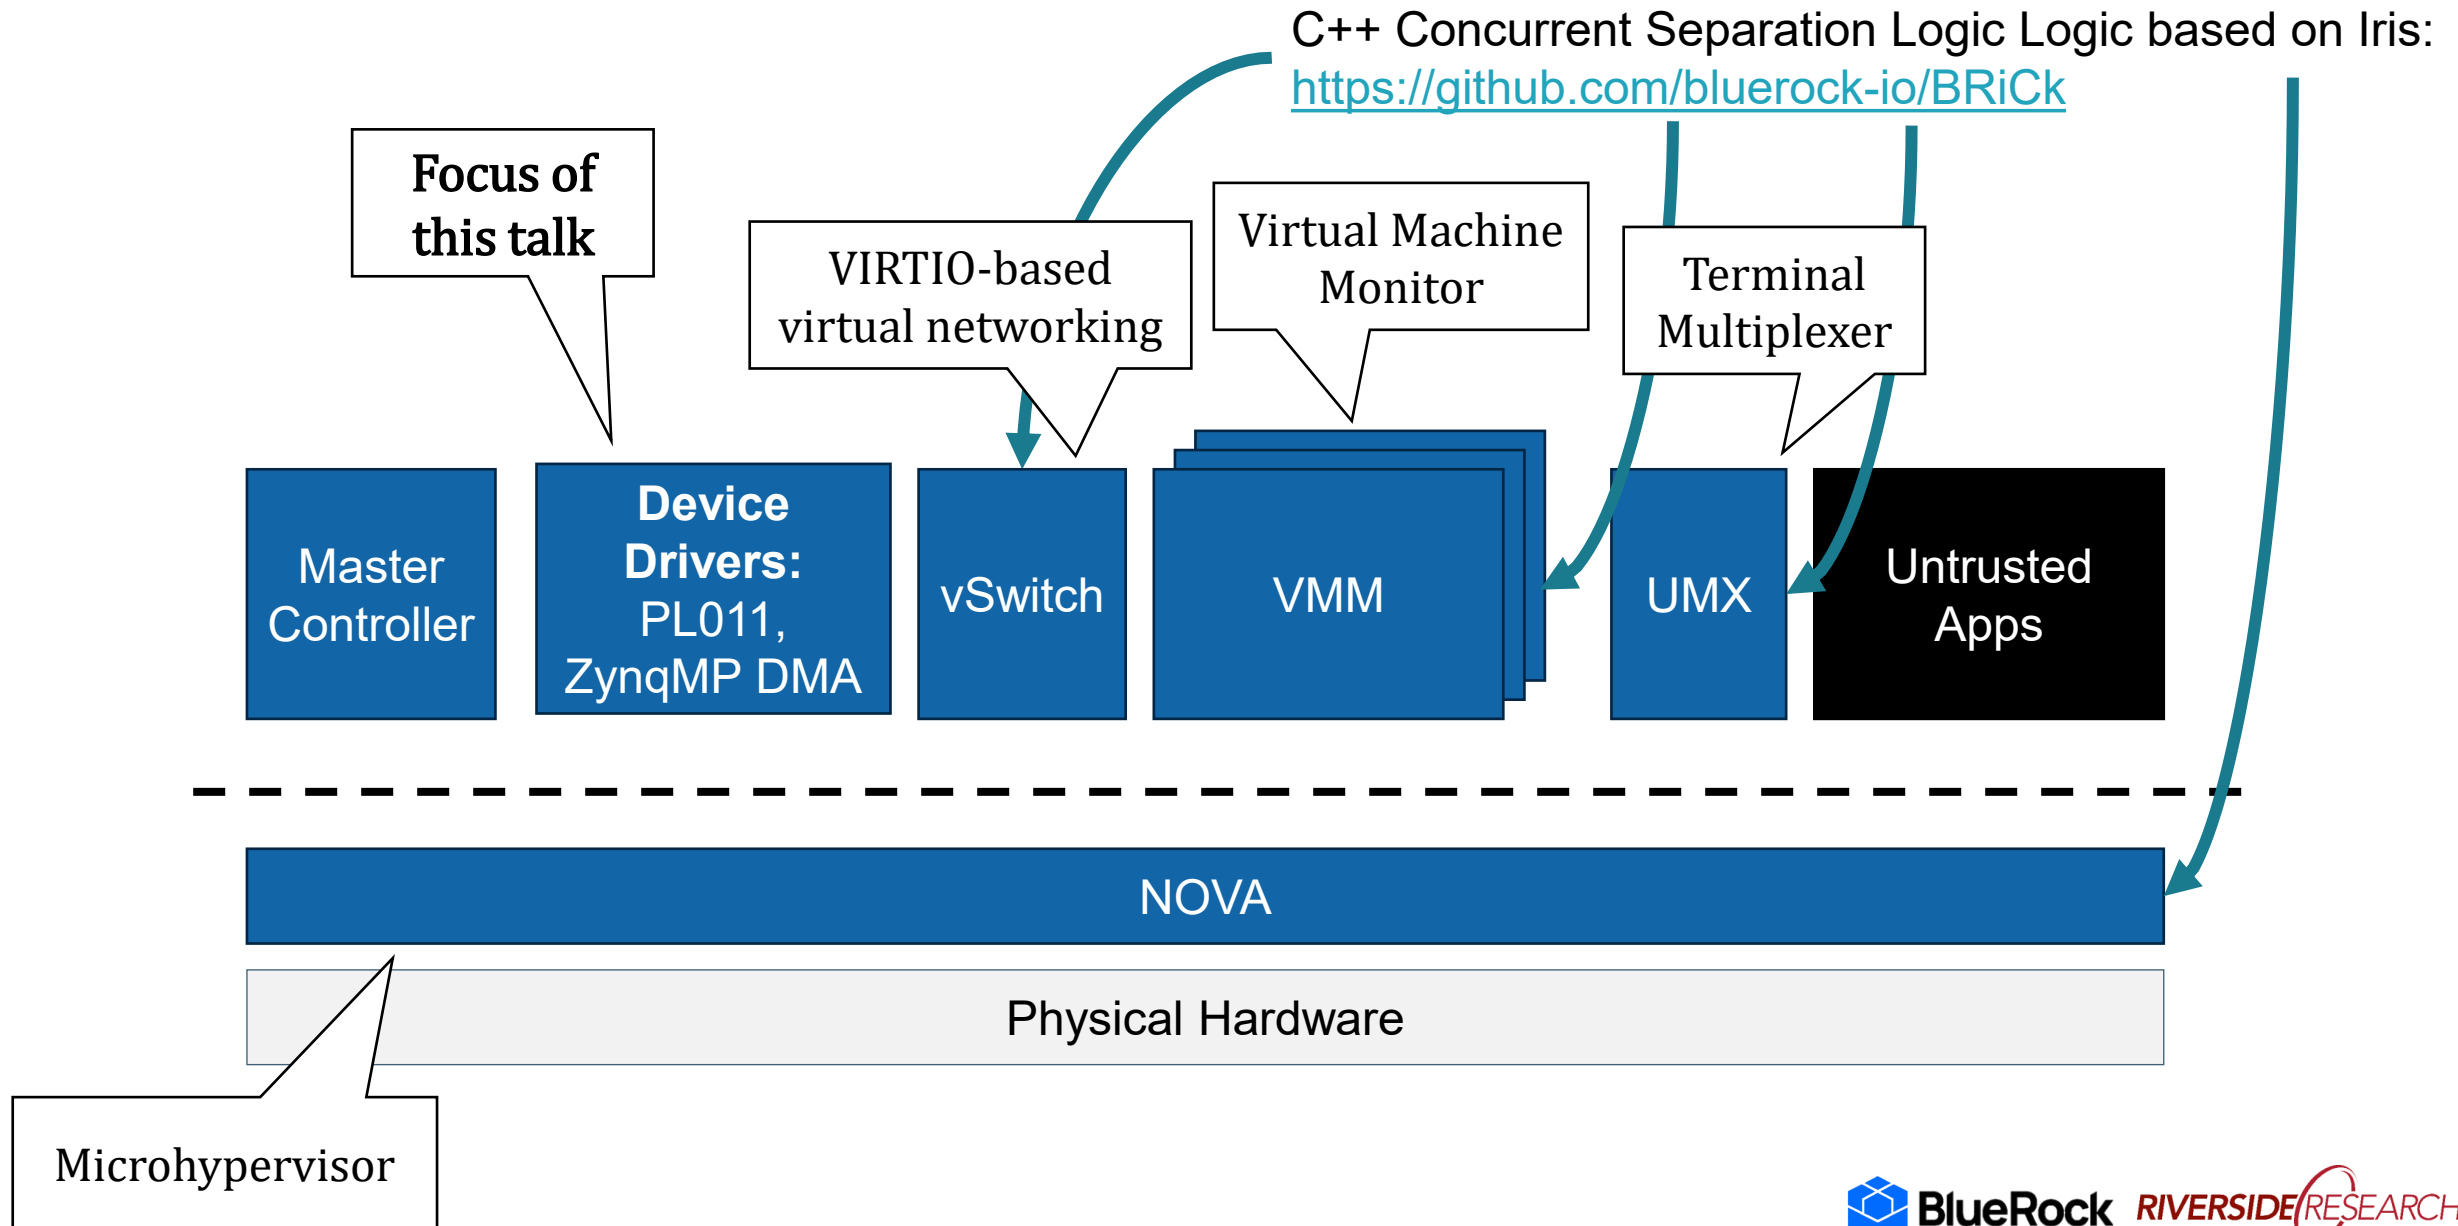
Verified ZynqMP DMA Driver in Concurrent Separation Logic

┌ **Gordon Stewart**
Riverside Research

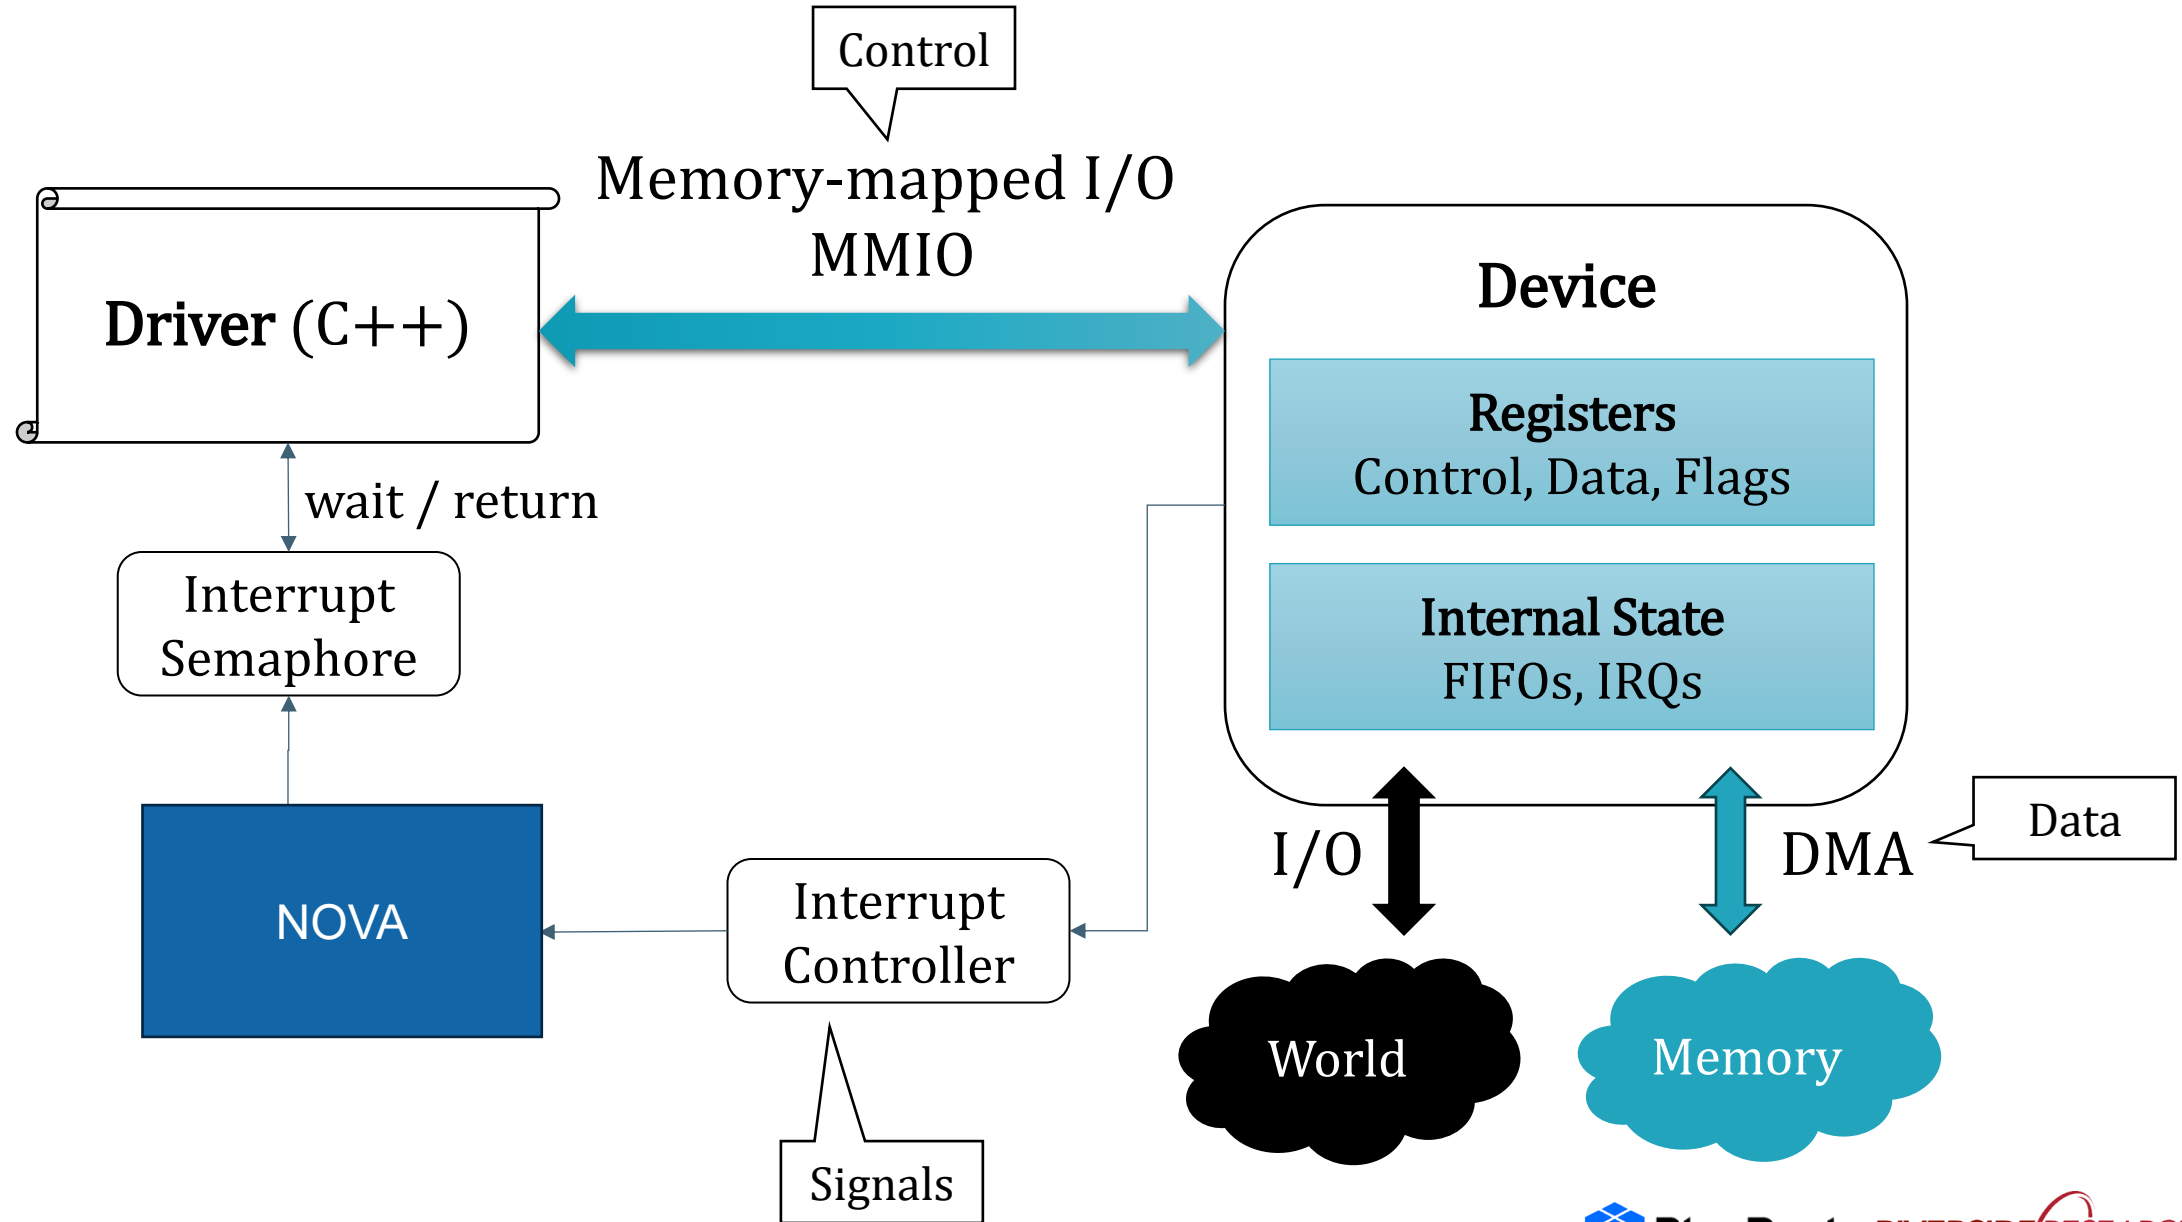
Gregory Malecha
BlueRock Security

seL4 Summit, Prague, Czechia
3 Sep 2025

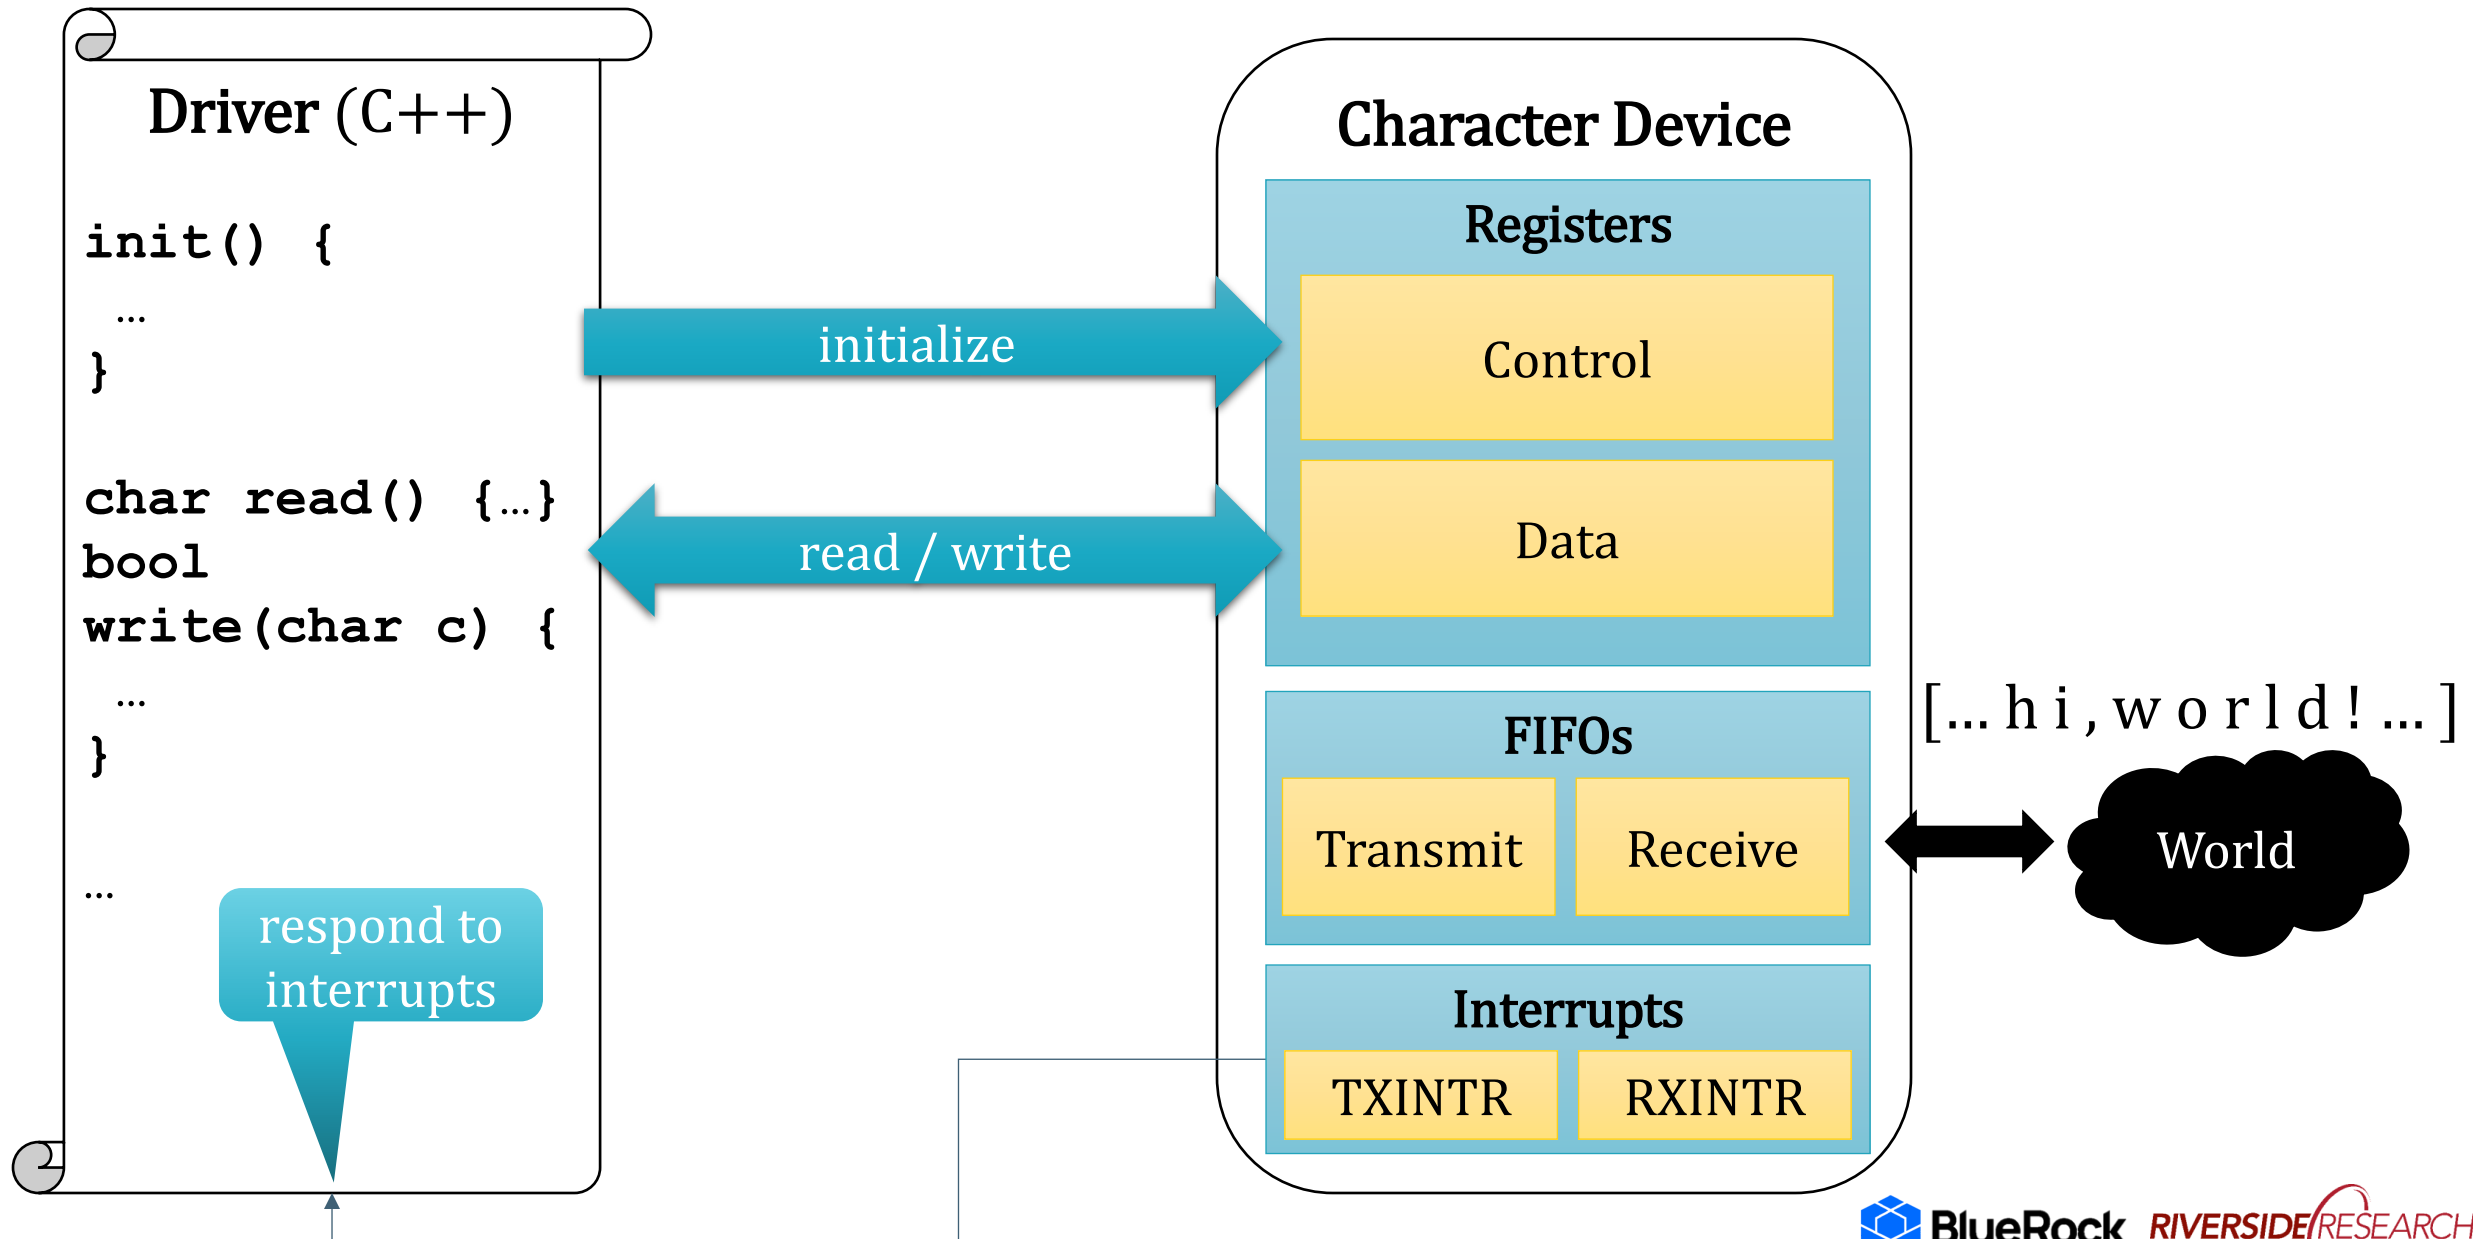
Operating Systems Verification at BlueRock Security



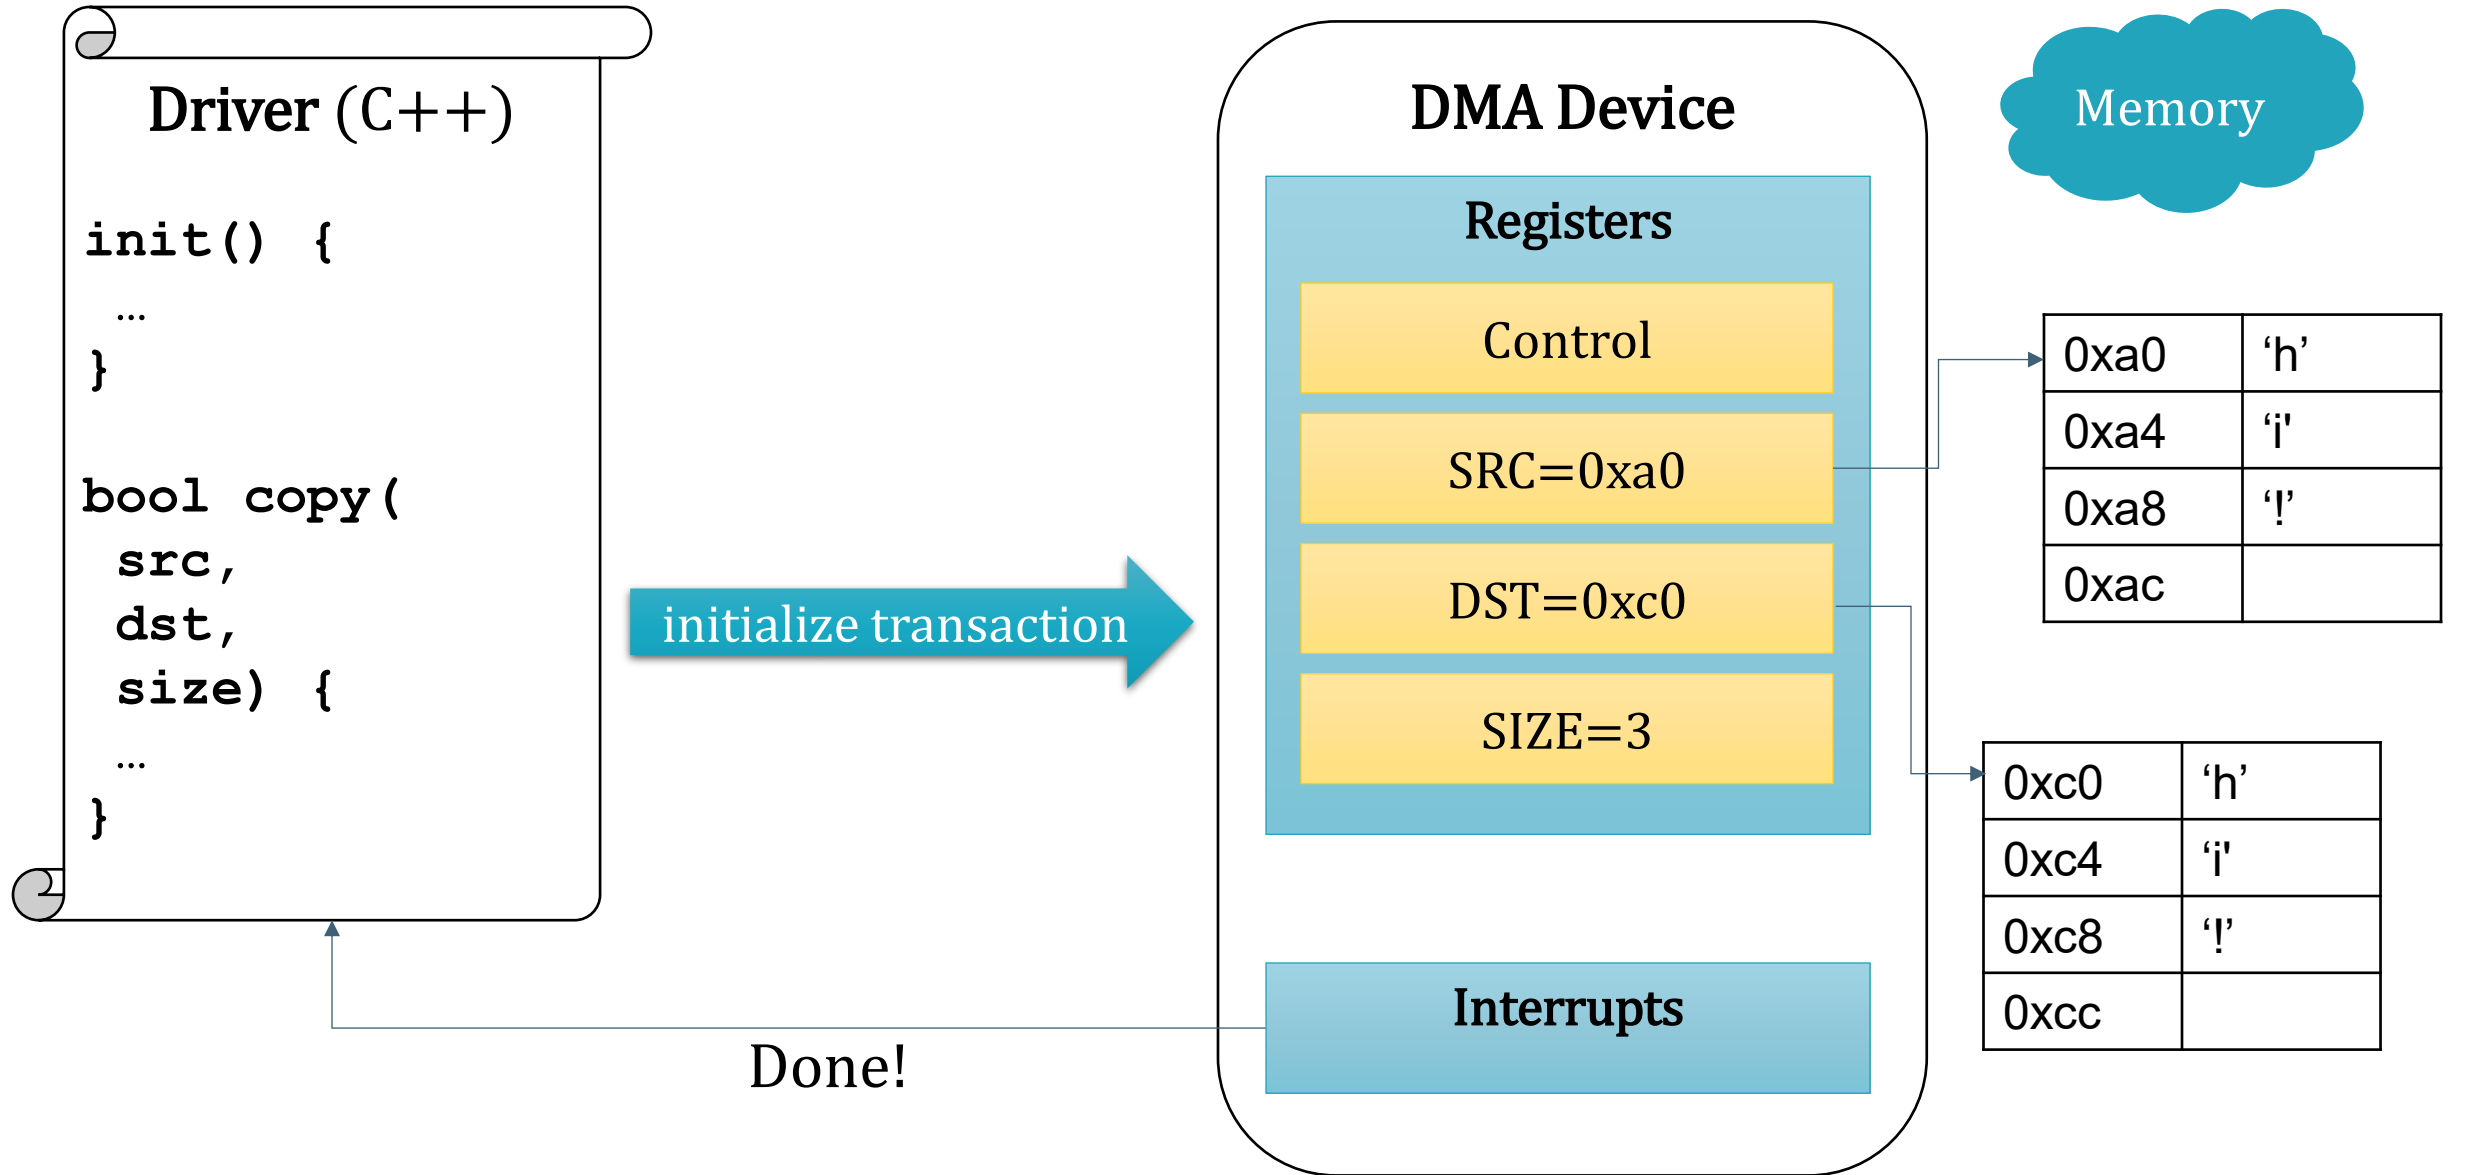
Device Drivers



Example #1: Character Device (UART)



Example #2: Direct-Memory Access (DMA, simple mode)



Outline

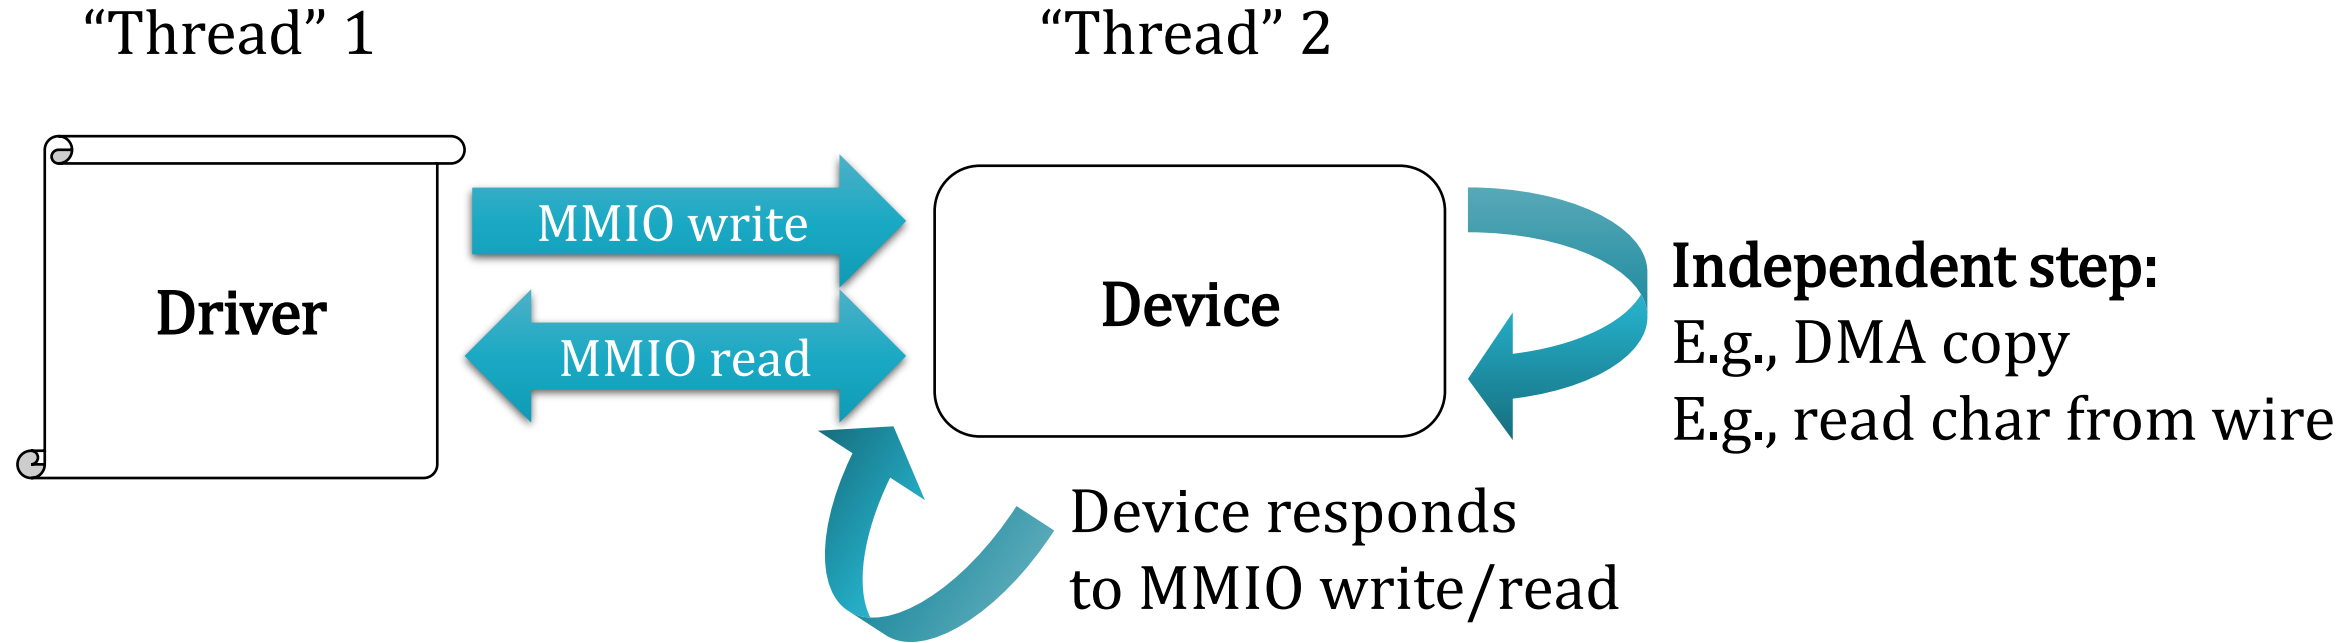
➔ Protocol-based verification by example:

PL011

ZynqMP DMA

seL4

Protocol-based Driver Verification

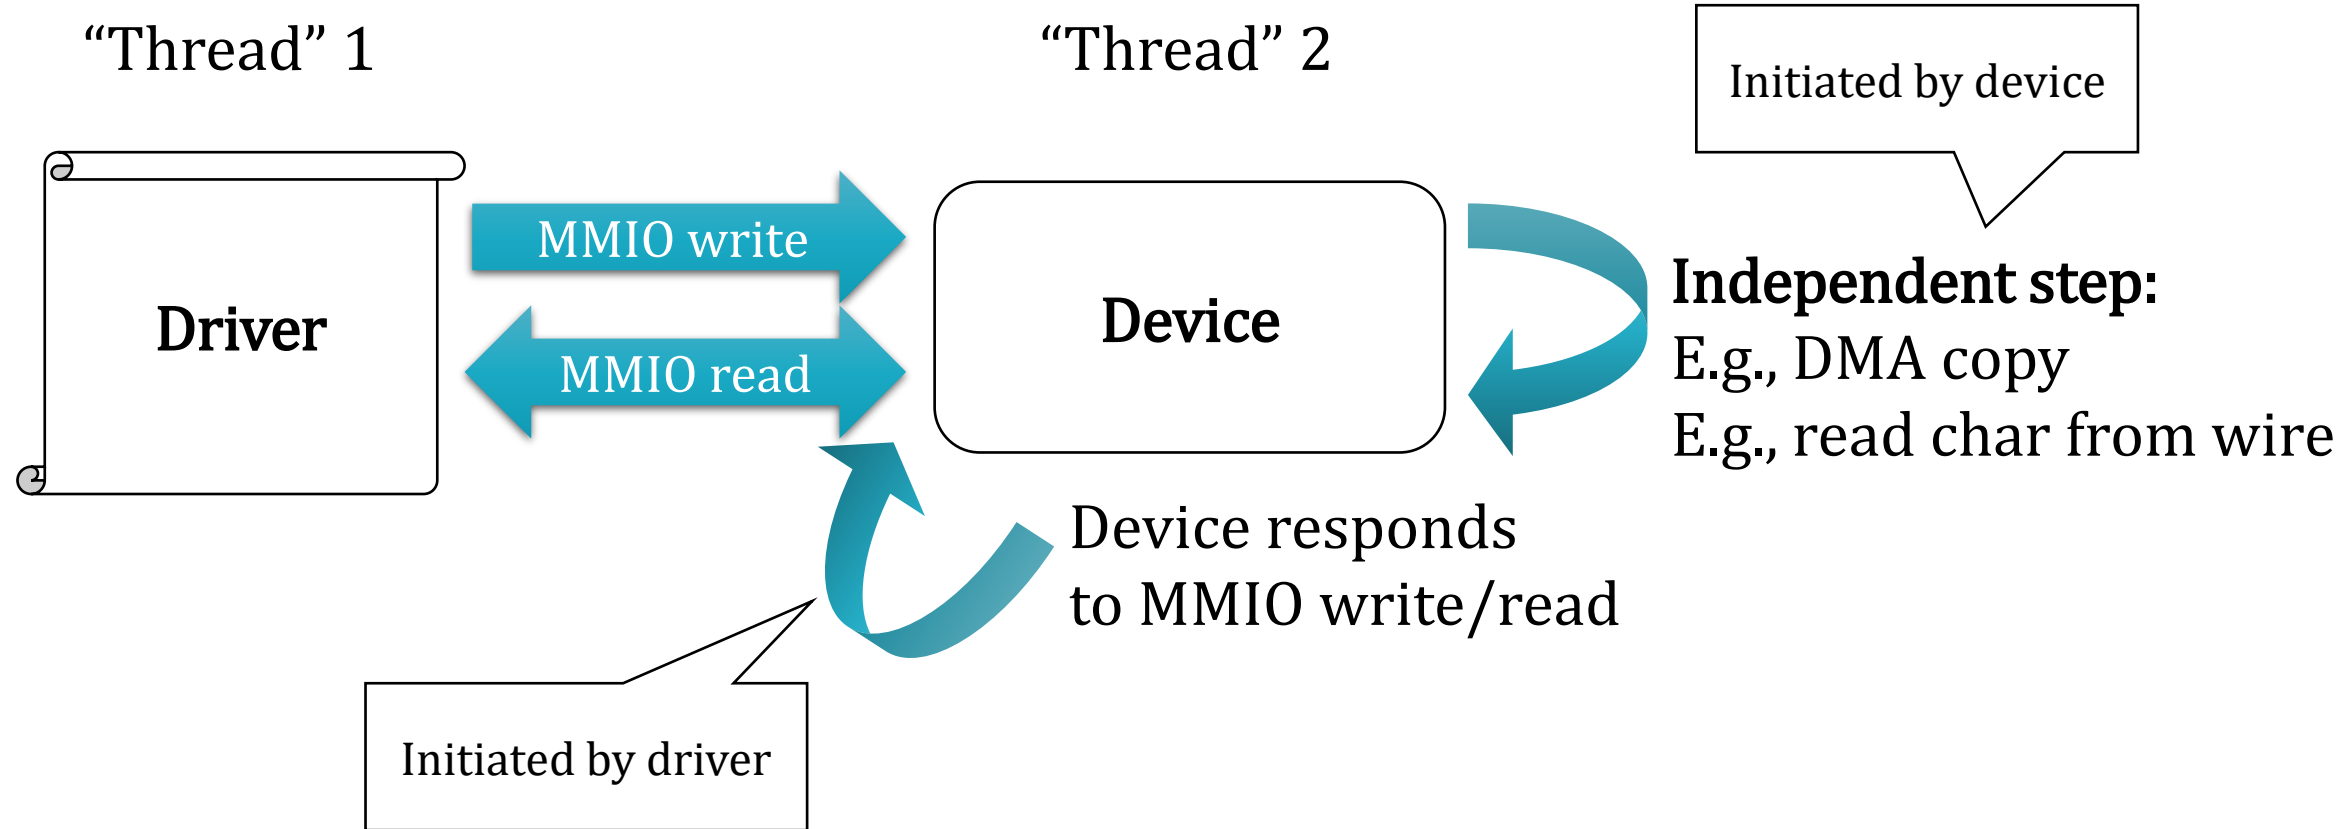


Device and driver run concurrently.

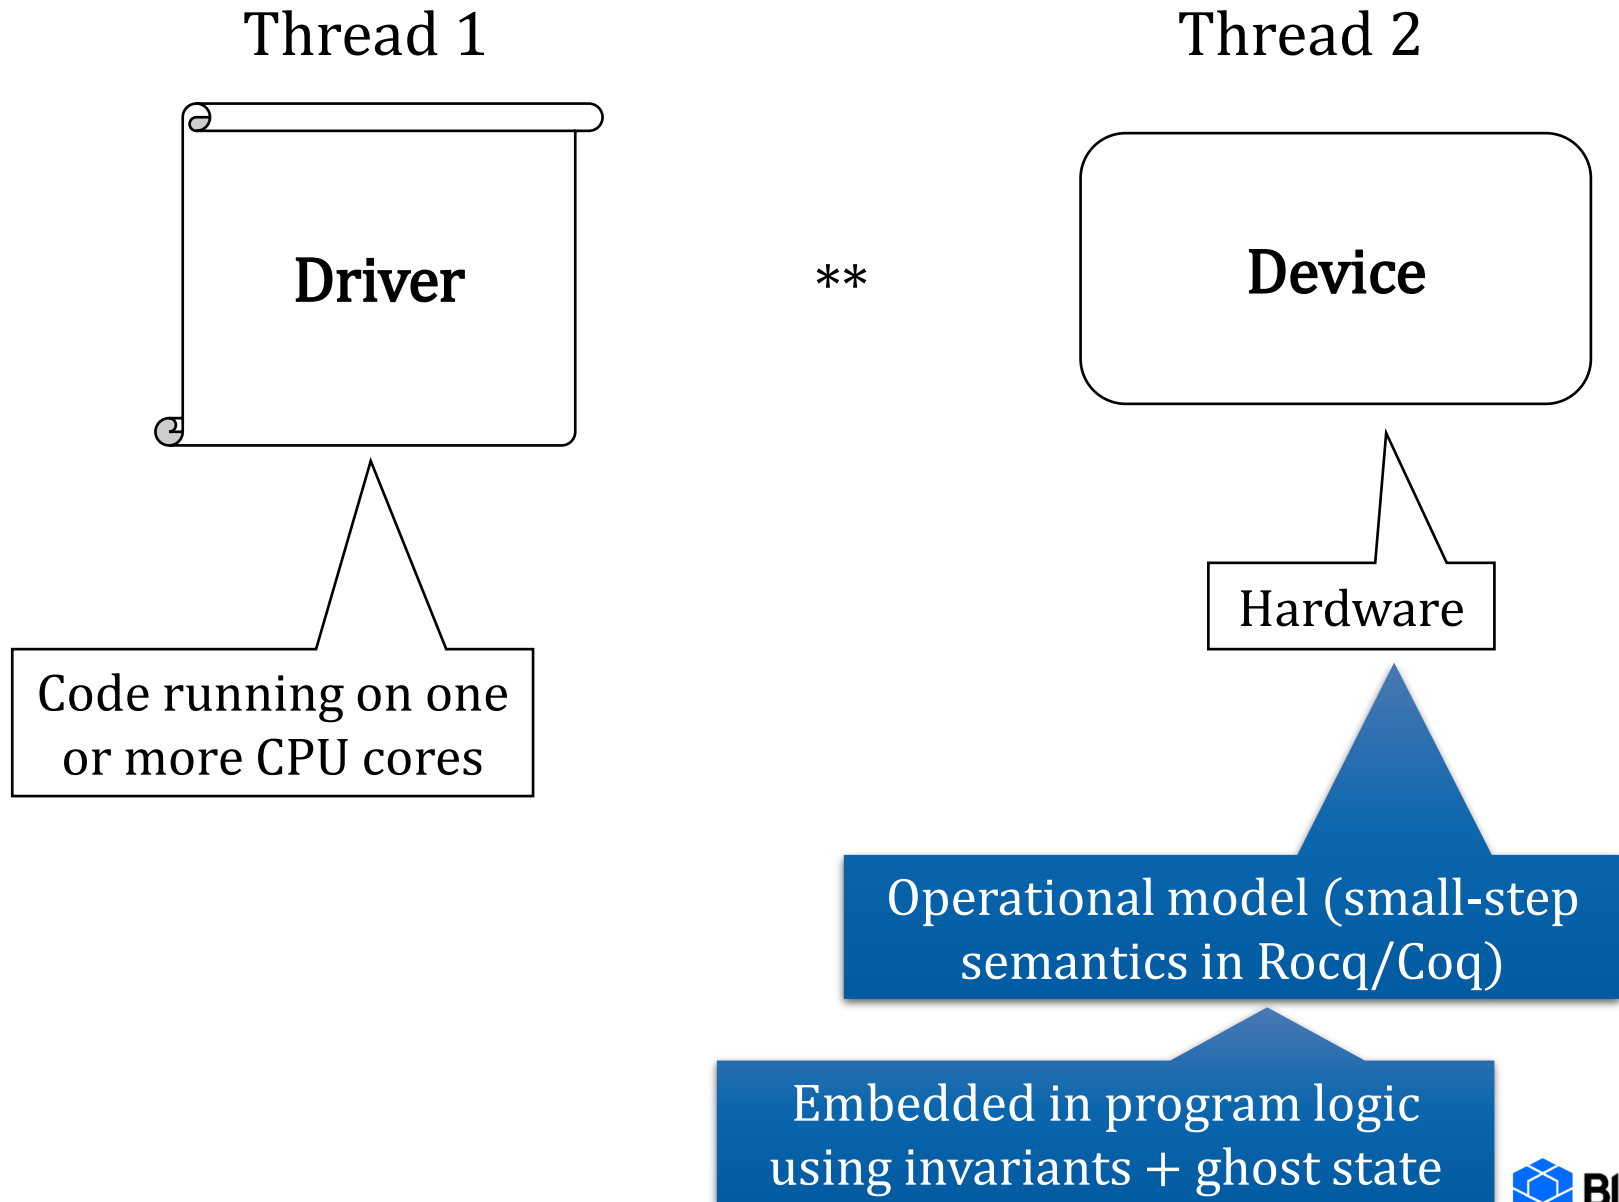
Treat each as a thread.

Interactions are on the shared state of the device.

Device and driver run concurrently



Treat driver and device as independent “threads”



Embedding Device States

Model state of a
UART device


```
(* Device State: *)  
#[local] Notation RxQ := (Queue.t N queue_size) (only parsing).  
#[local] Notation TxQ := (Queue.t N queue_size) (only parsing).  
Record State : Type :=  
  { rx : RxQ (* the receive buffer *)  
    ; tx : TxQ (* the transmit buffer *)  
    ; regs : @map Reg N (* the register state *)  
    ; irq_raised : @map Irq bool (* interrupts currently asserted *)  
    ; irqs : list IntAction }. (* pending irqs (not yet delivered) *)
```

Device


```
| WRITE_DR s c (pf : ~Queue.full s.(tx))  
  (_ : s .^ _uart_enabled) (_ : trimN8 c = c) :  
  let new_tx := Queue.enqueue c s.(tx) pf in  
  step s  
  (System (CpuWrite (natural_size_of DR) (offset_of DR) c))  
  (s &:      _tx .:= new_tx  
   &:      _txintr .:= false)
```

Example MMIO write step

Embedding Device States



Driver



Device

Iris authoritative camera

`State.frag γ st`

Driver view of state st

Ghost location containing device state

`State.auth γ st'`

Device view of state st'

Embedding Device States



Consistency:

`State.frag γ st ** State.auth γ st' -* [|st=st'|] ** ...`

Update:

`State.frag γ st ** State.auth γ st' -*
|=> State.frag γ new_st ** State.auth γ new_st`

Arbitrary new device state...

Embedding Device States

Invariant

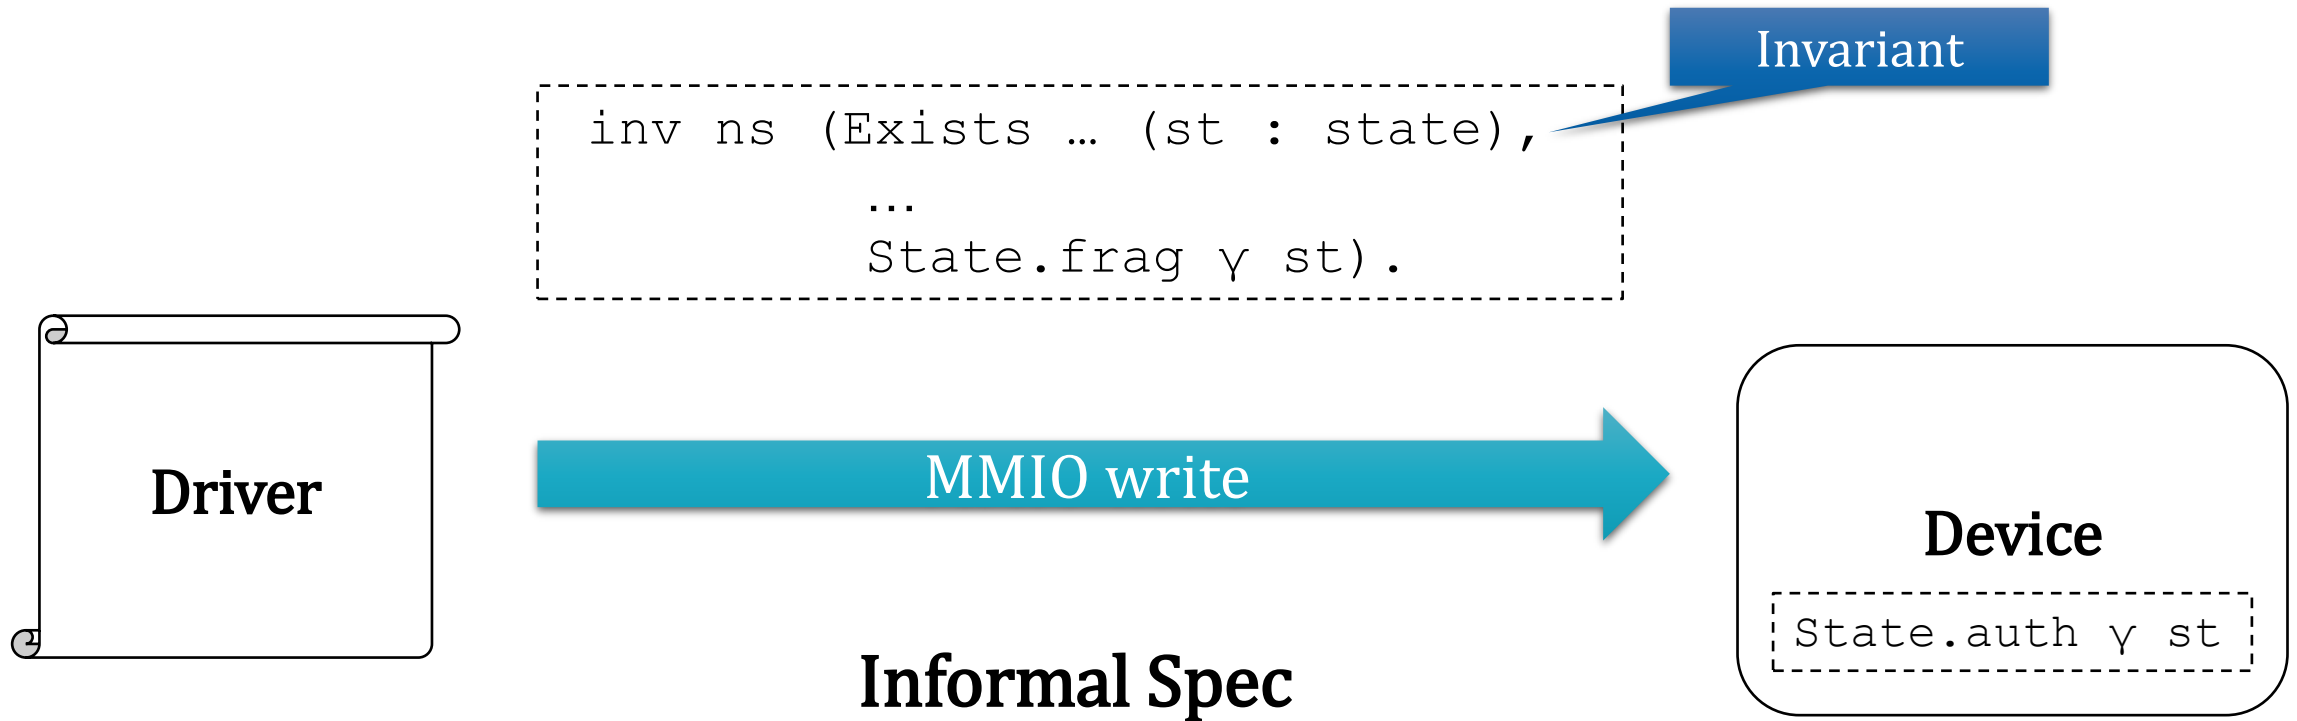
```
inv ns (Exists ... (st : state),  
    ...  
    State.frag γ st) .
```

Driver

Device

```
State.auth γ st
```

MMIO Operations



If `MMIO::write` is safe for current device state `st`,
Then `MMIO::write` updates device state to some `st'` reachable via a `CpuWrite` event.

Overall postcondition is some consequence Q .

MMIO write spec, formally

Definition `state_inv γ (ns : namespace) : mpred :=`
`inv ns (Exists ... (st : state),`
`...`
`State.frag γ.(_state) st).`

Assuming some device state
(will be constrained in a bit)

Definition `mmio_write_spec (sz : bitsize) : WpSpec_cpp :=`
`let ty := Tint sz Unsigned in`
`\with γdev offset Q ns`
`\arg{port} "port" (Vptr port)`
`\arg{val} "val" (Vn val) \prepost port |-> MMIOReg sz offset dev`
`\pre AU`
`<<V st, State.frag γdev st ** [| write_safe st sz offset val |]>>`
`@ protocol_mask ns, protocol_mask ns \ ↑stateinv_ns ns`
`<<E st',`
`[| dev_step st (System (CpuWrite (bitsN sz) offset val)) st' |] **`
`State.frag γdev st',`
`COMM Q>>`
`\post Q.`

... and it is safe to write to the
device register in this state ...

... then the device will update its state to `st'`,
in accord with its operational model.

Protocols to constrain device states

```
inv ns (Exists ... (st : state),  
    ...  
    State.frag γ st) .
```

Arbitrary state

Driver

```
State.frag γdev st **  
[| write_safe st sz offset val |]
```

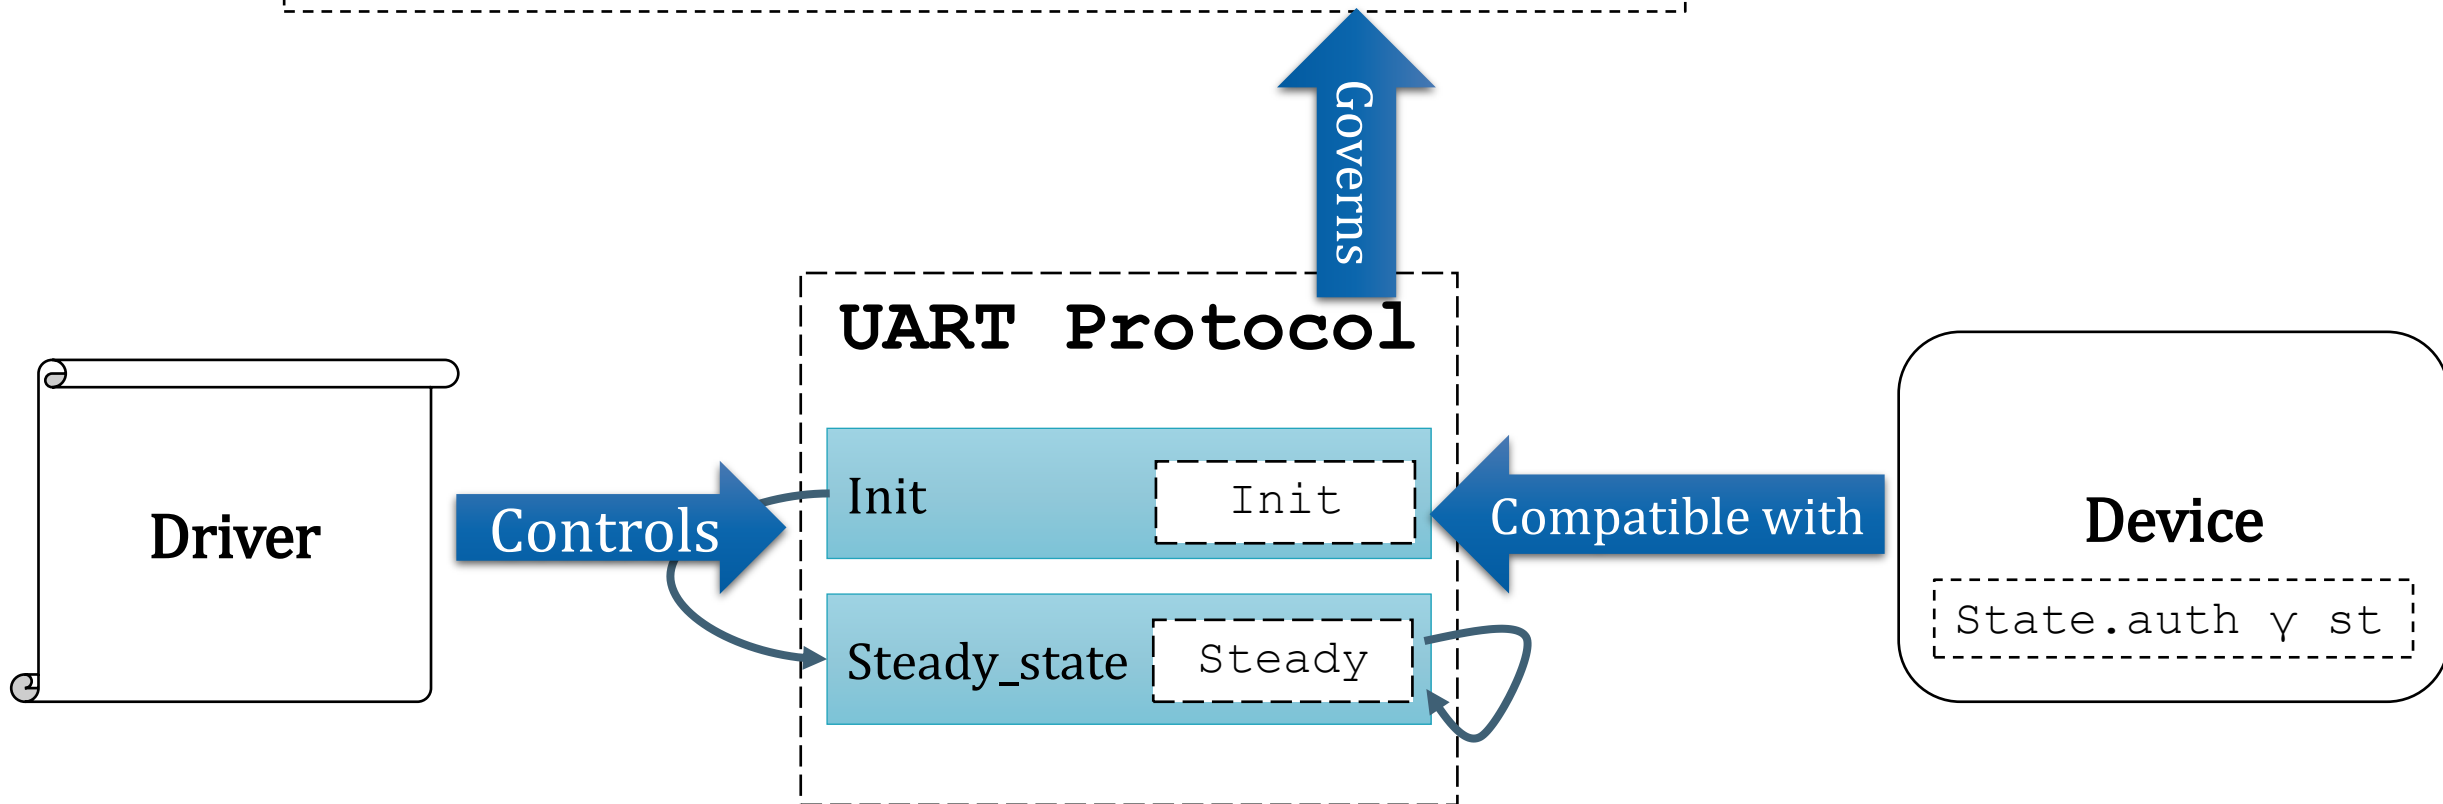
How to prove safety (and other deeper properties)?

Device

```
State.auth γ st
```


Protocols to constrain device states

```
inv ns (Exists (P : T) (st : state),  
  Proto.frag  $\gamma$ .(_proto) (1/2) P **  
  interp P st **  
  State.frag  $\gamma$ .(_state) st).
```

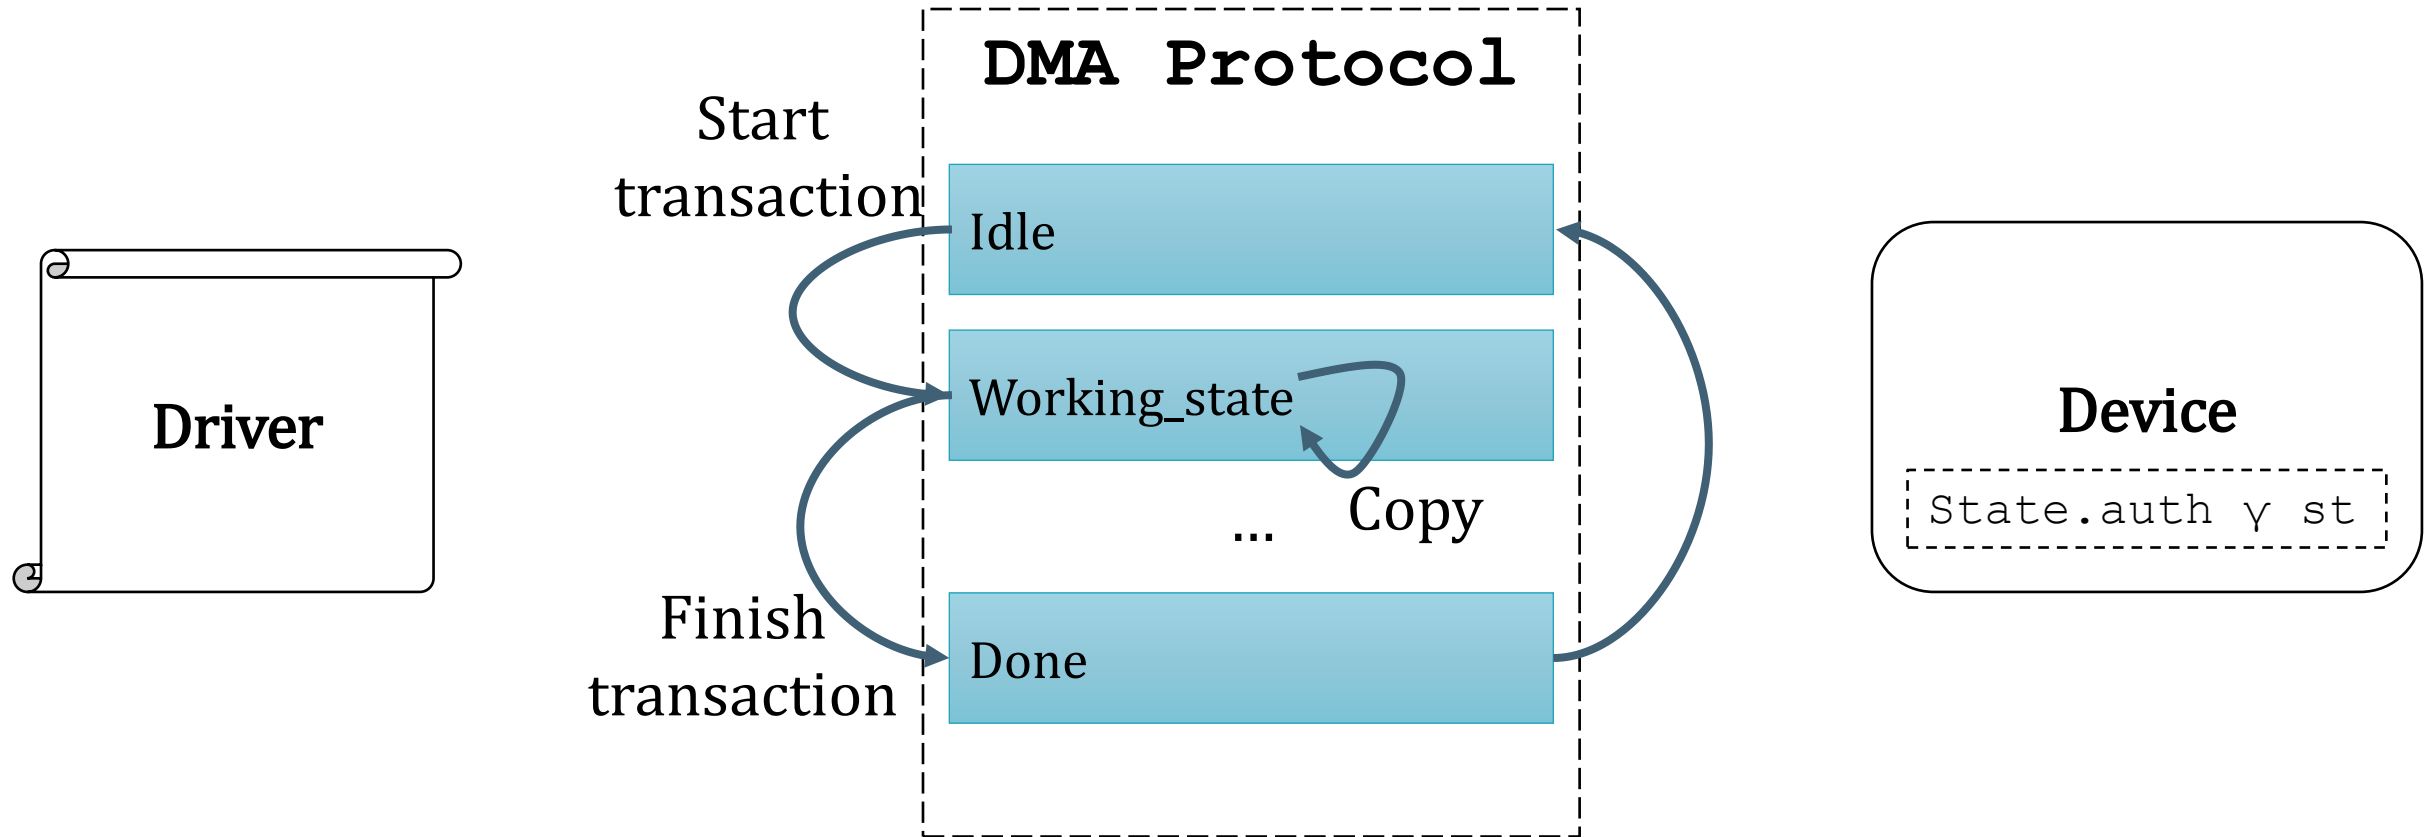


PL011: Interpretation of Steady_state

```
Definition Steady_state (γ : Chardev_model dev)
  : device_state -d> mpred0 := fun st =>
  [| Initialized st |] **
  ...
  (Exists 1 : leibniz0 (list N),
    [| 1 ++ fmap trimN8 (filter (negb ∘ receive_error)
      (Queue.list $ rxq st))
      = in_trace _ st |] **)
  own γ.(_intrace) (A:=excl_authR (leibniz0 (list N)))
    (excl_auth_auth 1)) **
  ...
```

Characters received so far, up to transmission errors and characters still in RX queue

ZynqMP DMA Protocol



Working_state

Definition wp_copy

```
(c : copy) (* Copy continuation *)
(sptr dptr : ptr) (* Ptr start of src/dest arrays *)
(l : list N) (* Data to be copied *)
(q : Qp) (* Ownership of src transferred to device *)
: mpred :=
  let s := c.(src_base) in
  let d := c.(dst_base) in
  ...
  let to_write := (sizen - num_written)%N in
  pinned_bytes size s sptr ** pinned_bytes size d dptr **
  [| c.(len) = N.of_nat (List.length l) |] **
  [| (s <= c.(src) /\ c.(src) <= s + sizen)%N |] **
  [| (d <= c.(dst) /\ c.(dst) <= d + sizen)%N |] **
  [| l = take (N.to_nat num_written) l ++
    c.(data) ++
    drop (N.to_nat num_read) l |] **
  [| l = map (fun n => (n `mod` 2 ^ 8)%N) l |] **
  sptr |-> bytesR q l **
  dptr |-> bytesR 1 (take (N.to_nat num_written) l) **
  dptr .[ T_uchar ! num_written ] |-> anyR (Tarray T_uchar to_write) 1.
```

Protocol phase corresponding to
“currently performing a copy”

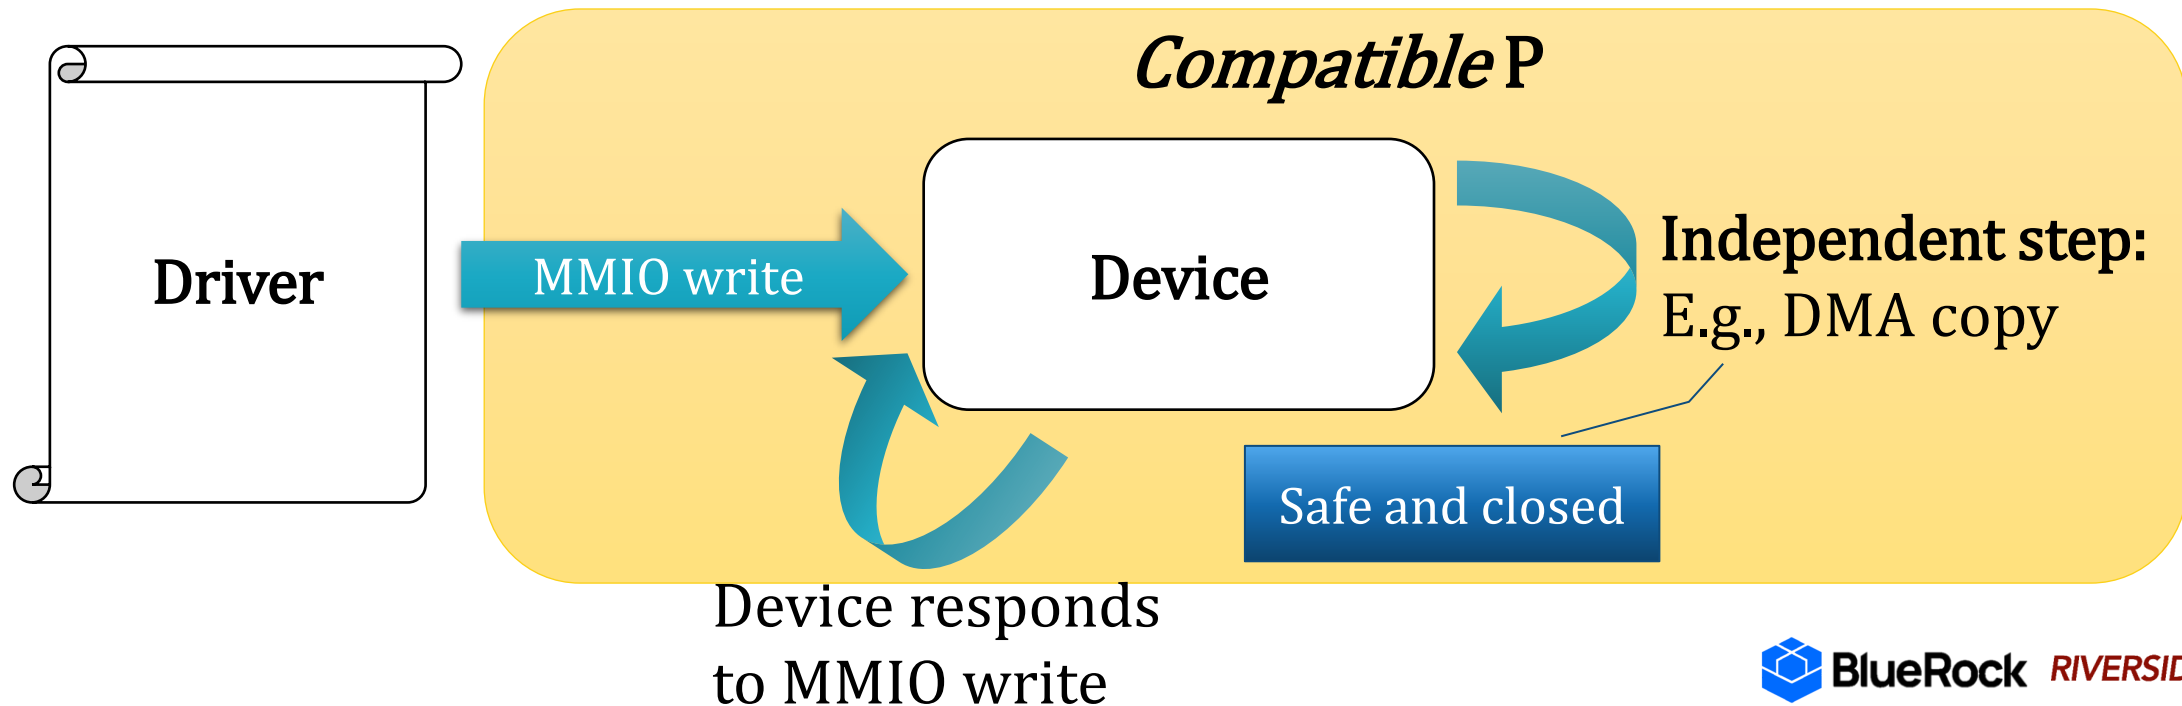
`l` is overall list of bytes to
be copied

Device is somewhere in the
middle of a copy

Protocols cannot block devices

Definition `state_inv γ (ns : namespace) : mpred :=`
`inv ns (Exists (P : T) (st : state),`
`Proto.frag γ.(_proto) (1/2) P **`
`interp P st **`
`State.frag γ.(_state) st).`

Arbitrary P could constrain the device state in ways that are *incompatible* with steps device takes independently of driver.



Protocols cannot block devices

Definition `state_inv` γ (ns : namespace) : mpred :=
inv ns (Exists (**P : T**) (**st : state**),
 Proto.frag γ ._proto) (1/2) **P** **
 interp **P** **st** **
 State.frag γ ._state) st).

Constrain the current protocol to
be **Compatible** with the device

Definition `compatible` γ (ns : namespace) :=
inv ns (Exists (P : T),
 Proto.auth γ ._proto) **P** **
 Compatible (interp P)).

`Compatible` (I : state -> mpred) : mpred :=
(\square (Forall st, I st -* wp_step st I))%I.

I is closed and safe for device steps

Weakest precondition predicate
transformer for device semantics

PL011 read_byte

Definition read_byte_spec (this : ptr) : WpSpec_cpp :=

\with γ q before

\arg{out} "out" (Vptr out)

\arg{timeout_tsc} "timeout_tsc" (Vn timeout_tsc)

\prepost this |-> chardev_drvR γ q

\pre In_trace γ before

\pre out |-> anyR T_uchar 1

\post{r}[Verrno r]

if decide (r = ENONE) then

Exists c, In_trace γ (before ++ [c]) ** out |-> uint8R 1 c

else In_trace γ before ** out |-> anyR T_uchar 1.

No error → Input trace extension

Simple-mode ZynqMP DMA



```
Definition memcpy_spec_ (this : ptr) : WithPrePost _ :=
  \with γ_proto γ γ_dma sptr dptr q l
  \arg{dest_pa} "dest_pa" (Vn dest_pa)
  \arg{source_pa} "source_pa" (Vn source_pa)
  \arg "len" (Vn (N.of_nat (List.length l)))
  \arg{timeout_tsc} "timeout_tsc" (Vn timeout_tsc)
  \require 0 < List.length l
  \require bound W32 Unsigned (Z.of_nat (List.length l))
  \prepost this |-> ZynqdmaR (Proto_idle γ) γ_proto γ_dma proto_full
  \prepost driver_regs γ
  \prepost pinned_bytes (List.length l) source_pa sptr
  \prepost pinned_bytes (List.length l) dest_pa dptr
  \prepost sptr |-> bytesR q l
  \pre dptr |-> anyR (Tarray (Tint char_bits Unsigned) (N.of_nat (length l))) 1
  \post{r}[Verrno r]
    if decide (r = ENONE)
      then dptr |-> bytesR 1 1
      else dptr |-> anyR (Tarray (Tint char_bits Unsigned) (N.of_nat (length l))) 1.
```

Device starts and ends in
Proto_idle state

Bytes 1 copied to
destination buffer

```
Definition memcpy_spec := [CHECK]
  SPECIFY (exact "_ZN10Zynqmp_dma7memcpyEmmmy") memcpy_spec_.
```


seL4 Connections


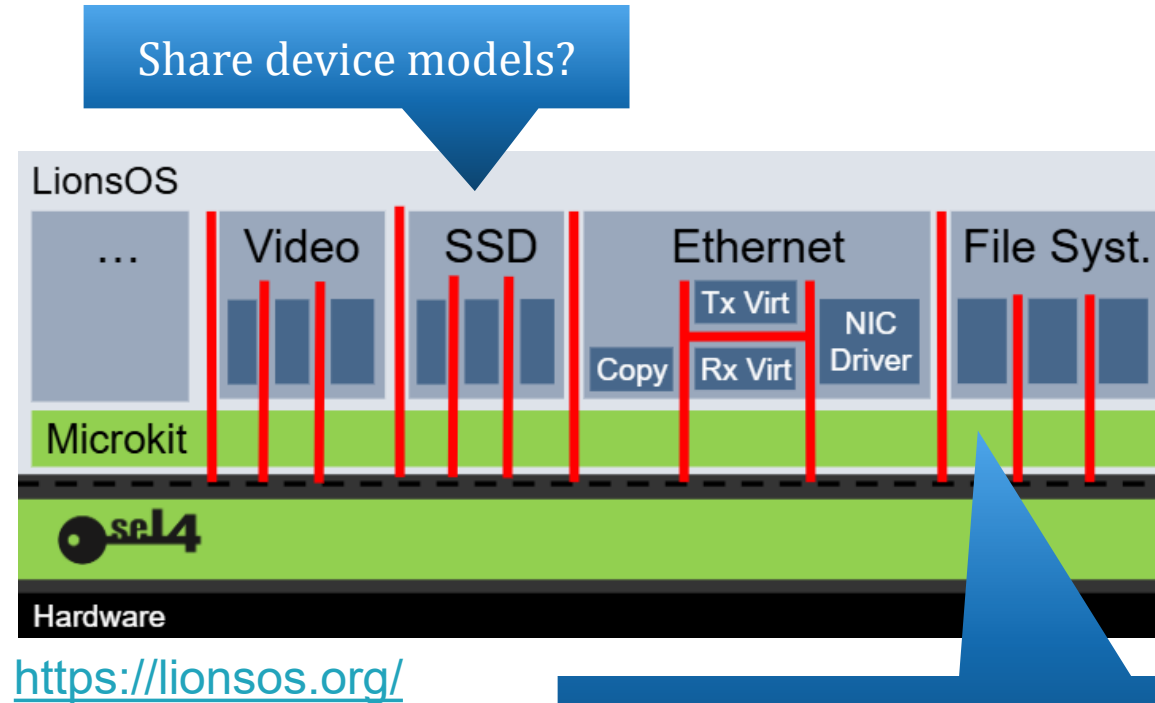


Figure 3.1: Device model.

<https://trustworthy.systems/projects/drivers/sddf-design.pdf>



Thanks!



Gordon Stewart

Riverside Research

gstewart@riversideresearch.org

Gregory Malecha

BlueRock Security

gregory@bluerock.io